

NAME: MARK LOPES

ROLL NO: 9913

BRANCH: COMPUTER A

EXP 6: NAIVEBAYES CLASSIFICATION

Using the [social network ads data](#) concerning the Gender, Age, and Estimated Salary of several users and based on these data we would classify each user whether they would purchase the insurance or not.

```
1 # Import libraries
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 from matplotlib.colors import ListedColormap
6 import seaborn as sns
7 from sklearn.preprocessing import LabelEncoder
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.model_selection import train_test_split
10 from sklearn.naive_bayes import GaussianNB
11 from sklearn import metrics
12 from sklearn.metrics import accuracy_score
13 from sklearn.metrics import classification_report
14 from sklearn.metrics import precision_recall_curve
15 from sklearn.metrics import confusion_matrix
16 from sklearn.metrics import f1_score
```

```
1 import pandas as pd
2 df = pd.read_csv('Social_Network_Ads.csv')
3 df
```

Output:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

```
[3] 1 df.head()
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
1 # Get required data
2 df.drop(columns = ['User ID'], inplace=True)
3 df.head()
```

	Gender	Age	EstimatedSalary	Purchased
0	Male	19	19000	0
1	Male	35	20000	0
2	Female	26	43000	0
3	Female	27	57000	0
4	Male	19	76000	0

```
1 df.describe()
```

	Age	EstimatedSalary	Purchased
count	400.000000	400.000000	400.000000
mean	37.655000	69742.500000	0.357500
std	10.482877	34096.960282	0.479864
min	18.000000	15000.000000	0.000000
25%	29.750000	43000.000000	0.000000
50%	37.000000	70000.000000	0.000000
75%	46.000000	88000.000000	1.000000
max	60.000000	150000.000000	1.000000

De



```
1 sns.distplot(df['EstimatedSalary'])
```




<ipython-input-29-cca866f85a80>:1: UserWarning:

'distplot' is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either 'displot' (a figure-level function with similar flexibility) or 'histplot' (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['EstimatedSalary'])  
<Axes: xlabel='EstimatedSalary', ylabel='Density'>
```



De

```
[7] 1 # Label encoding  
2 le = LabelEncoder()  
3 df['Gender'] = le.fit_transform(df['Gender'])
```

De



```
1 # Correlation matrix  
2 df.corr()  
3 sns.heatmap(df.corr())
```



<Axes: >



```
[10] 1 # Drop Gender column
      2 df.drop(columns=['Gender'], inplace=True)

[11] 1 # Split data into dependent/independent variables
      2 X = df.iloc[:, :-1].values
      3 y = df.iloc[:, -1].values

[12] 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = True)

[13] 1 # Scale dataset
      2 sc = StandardScaler()
      3 X_train = sc.fit_transform(X_train)
      4 X_test = sc.transform(X_test)

1 # Classifier
2 classifier = GaussianNB()
3 classifier.fit(X_train, y_train)
```

↕ GaussianNB ⓘ ⓘ
GaussianNB()

```
1 # Prediction
2 y_pred = classifier.predict(X_test)
3 print(np.concatenate((y_pred.reshape(len(y_pred), 1), y_test.reshape(len(y_test), 1)), 1))
```

↕

```
[1 1]
[0 0]
[1 1]
[1 1]
[0 0]
[0 1]
[0 0]
[1 1]
[1 0]
[1 1]
[1 0]
[0 0]
```

```
1 # Accuracy
2 accuracy_score(y_test, y_pred)
```

↕ 0.86

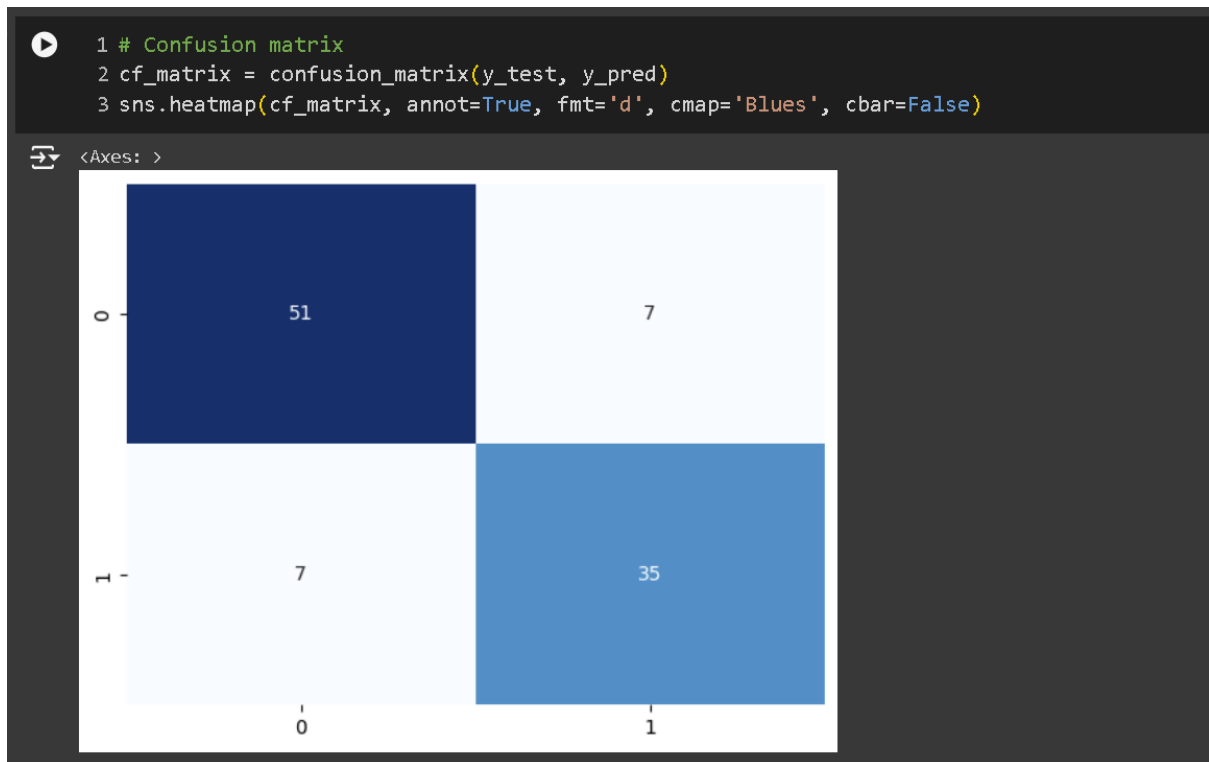
```
1 # Classification report
2 print(f'Classification Report: \n{classification_report(y_test, y_pred)}')
```

↕

Classification Report:				
	precision	recall	f1-score	support
0	0.88	0.88	0.88	58
1	0.83	0.83	0.83	42
accuracy			0.86	100
macro avg	0.86	0.86	0.86	100
weighted avg	0.86	0.86	0.86	100

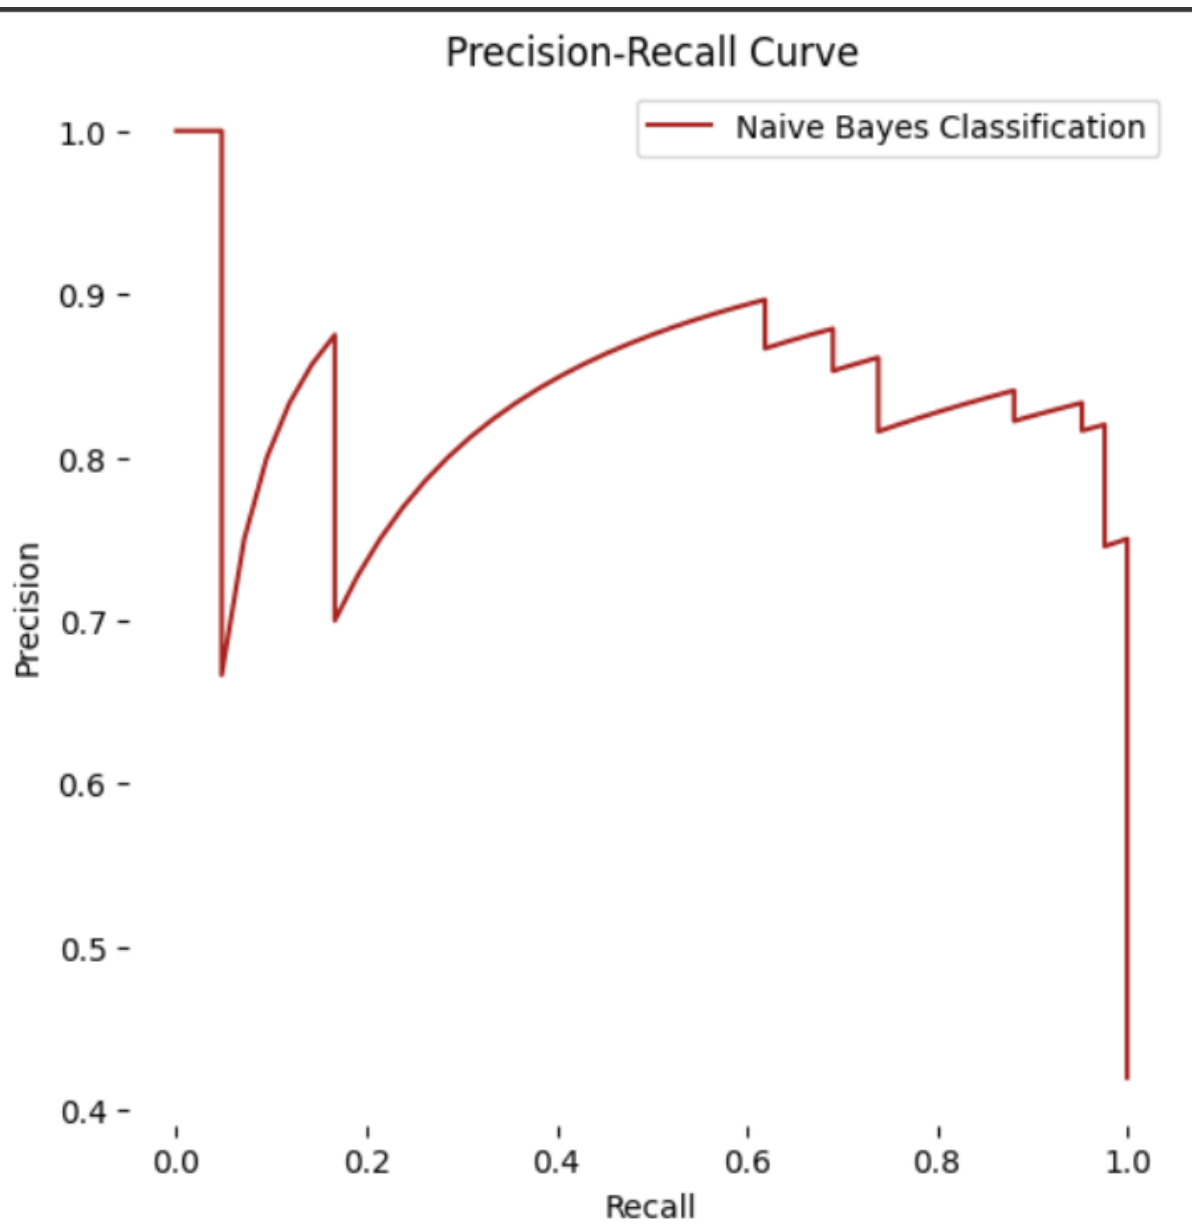
```
1 # F1 score
2 print(f"F1 Score : {f1_score(y_test, y_pred)}")
```

F1 Score : 0.8333333333333334




```
1 # Plot Precision-Recall Curve
2 y_pred_proba = classifier.predict_proba(X_test)[:,-1]
3 precision, recall, thresholds = precision_recall_curve(y_test, y_pred_proba)
4
5 fig, ax = plt.subplots(figsize=(6,6))
6 ax.plot(recall, precision, label='Naive Bayes Classification', color = 'firebrick')
7 ax.set_title('Precision-Recall Curve')
8 ax.set_xlabel('Recall')
9 ax.set_ylabel('Precision')
10 plt.box(False)
11 ax.legend();
```

Output:



```
1 # Plot AUC/ROC curve
2 y_pred_proba = classifier.predict_proba(X_test)[:,-1]
3 fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_proba)
4
5 fig, ax = plt.subplots(figsize=(6,6))
6 ax.plot(fpr, tpr, label='Naive Bayes Classification', color = 'firebrick')
7 ax.set_title('ROC Curve')
8 ax.set_xlabel('False Positive Rate')
9 ax.set_ylabel('True Positive Rate')
10 plt.box(False)
11 ax.legend();
```

Output:



FINAL OUTPUT SHOWING THE MODEL PREDICTING BASED ON AGE AND SALARY

```
[30] 1 # Testing the mdoel if user is likely to purchase the insurance or not.  
      2 # Predict purchase with Age(45) and Salary(97000)  
      3 print(classifier.predict(sc.transform([[45, 97000]])))
```

[1]

```
1 print(classifier.predict(sc.transform([[30, 9000]])))
```

[0]

```
[26] 1 print(classifier.predict(sc.transform([[60, 15000]])))
```

[1]